

Star Schema

Snowflakes

<<uses>>

RDBMS

Architecture

Data Warehouse

Decision Support System

Data Model

Properties

Fact

OLAP

Subject Oriented

Flexible Data Access

Non Volatile

Measure

Abstract OL

Backend

Time Stamp

Dimension

Attribute

Operation

MOLAP

ROLAP

Hybrid

• Hierarchical
• Let us perform fine grained aggregation

Sparsity problem

Joining Inefficient

We have a MOLAP backend wrt. to a ROLAP aggregation

Roll up

Drill Down

Slice And Dice

Drill Across

Pivoting

Aggregation = Group + Op

Roll up

Produces hypercubes

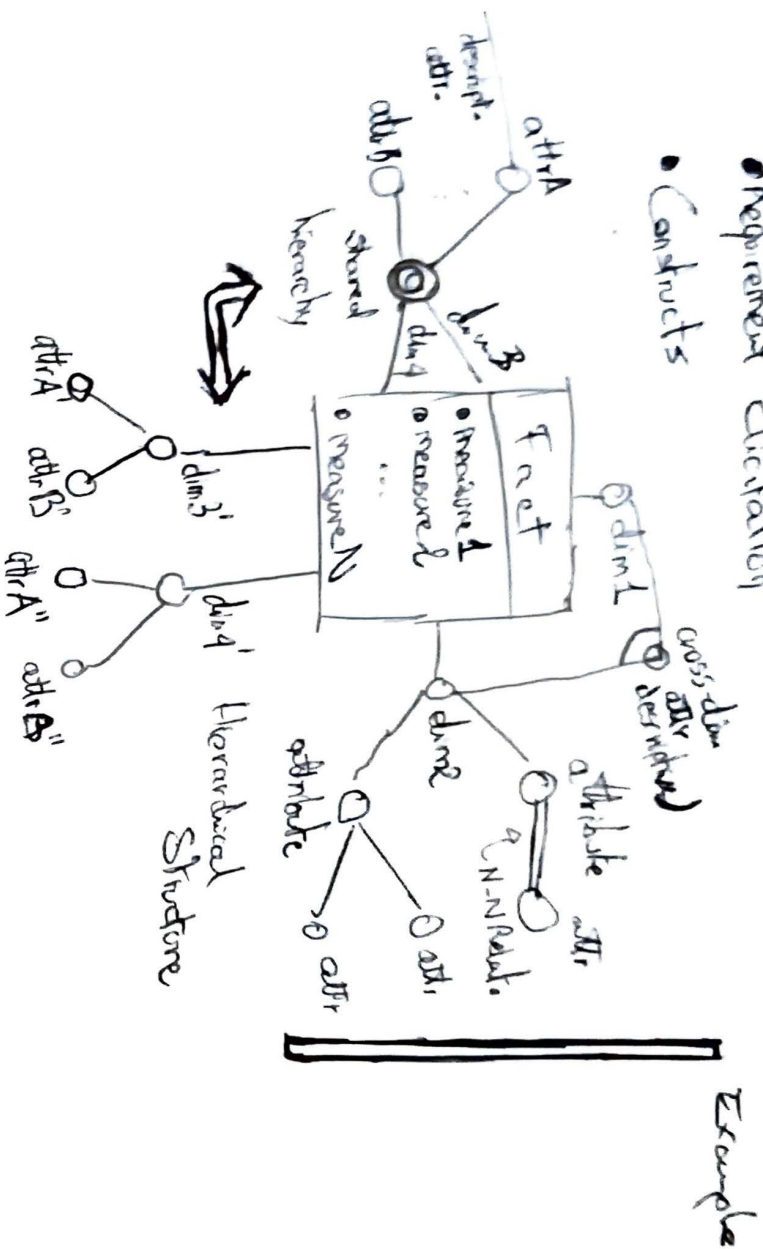
Compare away

"change the perspective of the query"


• No problems query
• Responsive

Design: Dimensional Fact Model

- Required Elucidation
- Constructs



- Measures are attribute we want to investigate (we want to assign them a value)
- Attributes are useful for aggregation
- Descriptive attributes are "additional values".
- Dimensions are just the root of an hierarchy.
- ... We can also have facts that need to be counted only ---



NOTE: An instance of a fact is ~~used~~ ^{an element} of the set of {dim3's ∪ attr's}

Homework Related Skill

ROLAP vs Star Schema vs Fact & Dim Tables vs Flatten the hierarchy
 The ^{solution} SnowFlakes vs Reduce Redundancy


* We don't need to "snowflake" the entire hierarchy (recursively)

Views

We can materialize aggregation exploring materialized view.

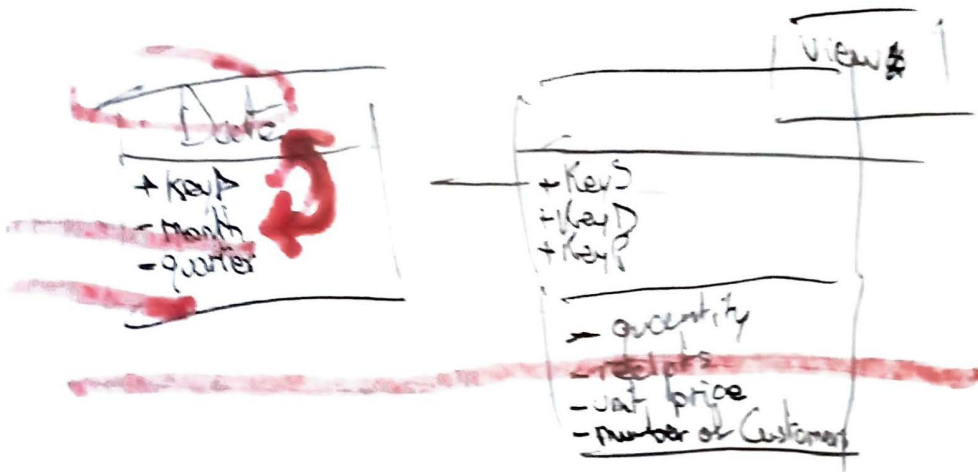

Notice that we can only store associative aggregate schema.

Example



We have only a surrogate for fun. no Derived measures

Slide = 106



Constellation Schema: Multiple stars (Remember the next 2 subsections)

Costellation Schema: Multiple stars (Remember the next 2 subsections)

A Typical architecture is 2-mic's architecture:

1. ~~10~~ Primary FACT (^)

2. Secondary FACT

NOTE:

- ~~State that~~ ~~exists~~ snowflakes ~~no~~ shared hierarchies ~~not~~ case of independence

We can start thinking of ... Quint Dimensions?

Bridge Table: N-to-N Relationship Bridge table: The manager POW

We use a ~~table~~ called "bridge table"

The query need to ~~be~~ consider weights to the arcs



No weight required Pitfall Example:

Materialized Aggregated View
association of assoc
create aggregate operation

Sale
+ book Key B
+ Key D
- number
- receipts

Book
+ Key B
- book
- genre

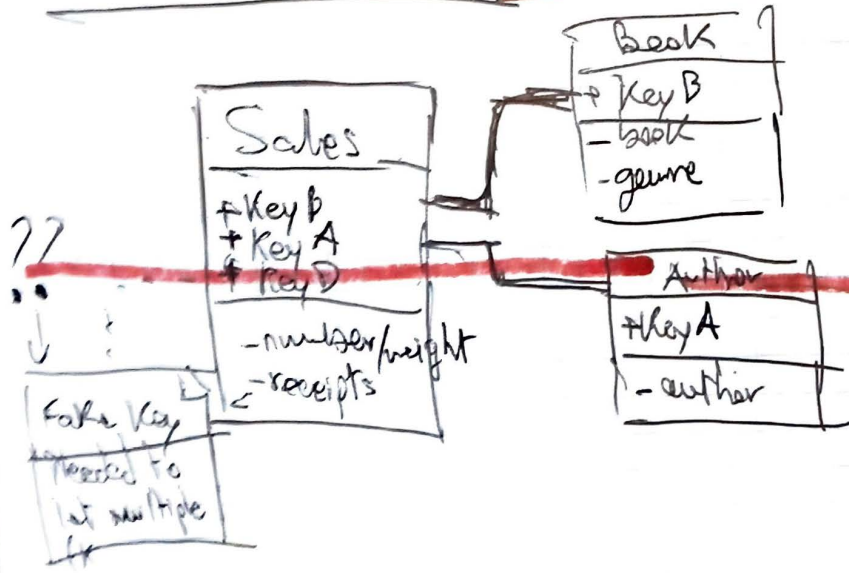
AUTHOR: author,
 Sum (SALES.receipts // index
 * BRIDGE.AUTH.weight)
 !
 AUTHOR: author
 Sum (SALES.number)
 NB:
 // Faculties of
 cannot do
 result 0

Author
+ Key A
- author

Bridge-Author
+ Key B
+ Key A
- <u>weight</u>

! IF ~~degenerate~~ non associative

Pushdown Solution: N-to-N relationship for non associativity, op.



Degenerate Dimensions

byte-len (attr-dims) \gg byte-len (synapse-key)

\Rightarrow ~~same~~
SUNX DIM.

~~Dimension takes degenerate~~
~~the attr bytes becomes dimensions~~
~~The attr byte caps exceeded is a difference~~
~~between get. degenerate in (unip) attr bytes~~

\Rightarrow

Less spin operation required (I'm taking it for granted)